

SEIS630:Database Mgmt Systems and Design

Benchmarking Cassandra Using Yahoo's Cloud Serving Benchmark (YCSB).

December 17th, 2021

- Shivali Dalmia



PROJECT WORKITEMS

- Difference between NoSQL and traditional SQL databases.
- Architecture of Cassandra NoSQL database.
- Analyzing performance of Cassandra database for different types of database operations using YCSB workloads.
- Experimental setup.
 - YCSB installation
 - Cassandra installation
 - Workload execution
- Comparison of throughput and latency for different types of workloads.



INSPIRATION

- To understand –
 - How distributed and non-distributed databases perform for read, insert and update operations.
 - How to select a database for different business use cases.
- To learn a new NoSQL database such as Cassandra.
- Key reference paper : Distributed NoSQL Data Stores: Performance Analysis and a Case Study by Abdeltawab Hendawi, Jayant Gupta, Jiayi Liu. [[Link](#)]



TOPICS FROM CLASS

- Difference between traditional SQL and NoSql databases.
- How CAP theorem applies to distributed and non distributed databases.
- Learned a new NoSql database Cassandra (e.g., architecture).
- How distributed databases perform on different database operations such as read, update, modify, write, delete.
- Role of performance metrics such as throughput and latency in analyzing different databases.



CHALLENGES

- Running a multi node cluster on a single windows machine.
- Setting up various multi node clusters with different configurations.
- Identifying appropriate sharding and replication factors.



FUTURE WORK

- Analyze the performance of Cassandra on:
 - Real world datasets.
 - Multi node cluster.
 - Nested queries.
- Compare and contrast Cassandra with different NoSQL databases.



ACKNOWLEDGEMENT

- Thank you, Prof. Abe Kazemzadeh, for your guidance.
- Kindly reach me out for any questions and concerns.
- Please refer to the final draft for more details.

